

## Semaphores and Mutexes

The POSIX thread library contains functions for working with semaphores and mutexes. There is much more to say than what is mentioned here. A good place to find more information is [linux.die.net](http://linux.die.net).

The functions should all be compiled and linked with -pthread.

### What is a semaphore in LINUX?

(Library: `#include <semaphore.h>`)

A semaphore is fundamentally an integer whose value is never allowed to fall below 0. There are two operations on a semaphore: wait and post. The post operation increments the semaphore by 1, and the wait operation does the following: If the semaphore has a value > 0, the semaphore is decremented by 1. If the semaphore has value 0, the caller will be blocked (busy-waiting or more likely on a queue) until the semaphore has a value larger than 0, and then it is decremented by 1. We declare a semaphore as:

**sem\_t sem;**

where sem\_t is a typedef defined in a header file as (apparently) a kind of unsigned char.

An example of this might be that we have a set of N interchangeable resources. We start with semaphore S = N. We use a resource, so there are now N-1 available (wait), and we return it when we are done (post). If the semaphore has value 0, there are no resources available, and we have to wait (until someone does a post).

Semaphores are thus used to coordinate concurrent processes.

This is what some people call a "counted semaphore". There is a similar notion called a "binary semaphore" which is limited to the values 0 and 1.

A semaphore may be named or unnamed. These notes assume we are using named semaphores.

### Semaphore Functions in C

#### 1. `int sem_init(sem_t * sem, int pshared, unsigned int value);`

Purpose: This initializes the semaphore \*sem. The initial value of the semaphore will be value. If pshared is 0, the semaphore is shared among all threads of a process (and hence need to be visible to all of them such as a global variable). If pshared is not zero, the semaphore is shared but should be in shared memory.

Notes:

- On success, the return value is 0, and on failure, the return value is -1.
- An attempt to initialize a semaphore that has already been initialized results in undefined behavior

#### 2. `int sem_wait(sem_t * sem);`

Purpose: This implements the wait function described above on the semaphore \*sem.

Notes:

- Here sem\_t is a typedef defined in the header file as (apparently) some variety of integer.
- On success, the return value is 0, and on failure, the return value is -1 (and the value of the semaphore is unchanged).
- There are related functions sem\_trywait() and sem\_timedwait().

### **3. int sem\_post(sem\_t \* sem);**

Purpose: This implements the post function described above on the semaphore \*sem.

Note: On success, the return value is 0, and on failure, the return value is -1 (and the value of the semaphore is unchanged).

### **4. int sem\_destroy(sem\_t \* sem);**

Prototype: int sem\_destroy(sem\_t \* sem);

Purpose: This destroys the semaphore \*sem, so \*sem becomes uninitialized.

Notes:

- On success, the return value is 0, and on failure, the return value is -1.
- Destroying a semaphore on which other processes or threads are waiting (using sem\_wait()) or destroying an uninitialized semaphore will produce undefined results.

---

## **What is a mutex in LINUX?**

(Library: #include <pthread.h>)

A mutex (named for "mutual exclusion") is a binary semaphore with an ownership restriction: it can be unlocked (the post operation) only by whoever locked it (the wait operation). Thus a mutex offers a somewhat stronger protection than an ordinary semaphore.

We declare a mutex as:

**pthread\_mutex\_t mutex;**

## **mutex Functions in C**

### **1. int pthread\_mutex\_init(pthread\_mutex\_t \* restrict mutex, const pthread\_mutexattr\_t \* restrict attr);**

Purpose: This initializes \*mutex with the attributes specified by attr. If attr is NULL, a default set of attributes is used. The initial state of \*mutex will be "initialized and unlocked".

Notes:

- If we attempt to initialize a mutex already initialized, the result is undefined.
- On success, the return value is 0, and on failure, the return value is a nonzero value indicating the type of error.

- In the prototype, the keyword restrict (part of the C99 standard) means that this pointer will be the only pointer to the object.

## 2. **int pthread\_mutex\_destroy(pthread\_mutex\_t \* restrict mutex);**

Purpose: This destroys the mutex object \*mutex, so \*mutex becomes uninitialized.

Notes:

- It is safe to destroy an unlocked mutex but not a locked mutex.
- The object \*mutex could be reused, i.e., reinitialized.
- On success, the return value is 0, and on failure, the return value is a nonzero value indicating the type of error.

## 3. **int pthread\_mutex\_lock(pthread\_mutex\_t \* mutex);**

Purpose: This locks \*mutex. If necessary, the caller is blocked until \*mutex is unlocked (by someone else) and then \*mutex is locked. When the function call ends, \*mutex will be in a locked state.

Notes:

- Suppose we try to relock a locked mutex. Depending on the attributes of the mutex, we may have an error, or a count may be kept of how many times the caller has locked the same mutex (and thus will have to unlock it the same number of times).
- On success, the return value is 0, and on failure, the return value is a nonzero value indicating the type of error.

## 4. **int pthread\_mutex\_unlock(pthread\_mutex\_t \* mutex);**

Purpose: This unlocks \*mutex.

Notes:

- Suppose we try to unlock an unlocked mutex. Depending on the attributes of the mutex, we may have an error.
- On success, the return value is 0, and on failure, the return value is a nonzero value indicating the type of error.

# Lab Tasks

## Task1:

Solve the producer and consumer problem with inter thread communication (join(), wait(), sleep() etc.) modifying the given C code.

```
#include <pthread.h>
#include <stdio.h>
#include <string.h>
#define MAX 10 //producers and consumers can produce and consume upto MAX
#define BUFLen 6
#define NUMTHREAD 2    /* number of threads */

void * consumer(int *id);
void * producer(int *id);

char buffer[BUFLen];
char source[BUFLen]; //from this array producer will store it's production into buffer
```

```
int pCount = 0;
int cCount = 0;
int buflen;
```

```
//initializing pthread mutex and condition variables
pthread_mutex_t count_mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t nonEmpty = PTHREAD_COND_INITIALIZER;
pthread_cond_t full = PTHREAD_COND_INITIALIZER;
int thread_id[NUMTHREAD] = {0,1};
int i = 0;
int j = 0;
```

```
main()
{
    int i;
    /* define the type to be pthread */
    pthread_t thread[NUMTHREAD];

    strcpy(source,"abcdef");
    buflen = strlen(source);
    /* create 2 threads*/
    /* create one consumer and one producer */
    /* define the properties of multi threads for both threads */
    //Write Code Here
```

```
}
```

```
void * producer(int *id)
{
    /*
        1. Producer stores the values in the buffer (Here copies values from source[] to buffer[]).
        2. Use mutex and thread communication (wait(), sleep() etc.) for critical section.
        3. Print which producer is storing which values using which thread inside the critical
        section.
        4. Producer can produce upto MAX
    */
    //Write code here
}
```

```
void * consumer(int *id)
```

```

{
    /*
    1. Consumer takes out the value from the buffer and makes it empty.
    2. Use mutex and thread communication (wait(), sleep() etc.) for critical section
    3. Print which consumer is taking which values using which thread inside the critical
    section.
    4. Consumer can consume upto MAX
    */
    //Write code here
}

```

### Example Output:

```

0 produced a by Thread 0
1 produced b by Thread 0
0 consumed by Thread 1
2 produced c by Thread 0
1 consumed b by Thread 1
3 produced d by Thread 0
4 produced e by Thread 0
2 consumed c by Thread 1
5 produced f by Thread 0
6 produced a by Thread 0
3 consumed d by Thread 1
7 produced b by Thread 0
8 produced c by Thread 0
9 produced d by Thread 0
4 consumed e by Thread 1

```

### Task 2:

The task is similar to the producer–consumer problem discussed in the class. The farmer and Shopowner share a fixed-size buffer named **warehouse** used as a queue. The farmer's job is to harvest crops(Rice=R, Wheat=W, Potato=P, Sugarcane=S, Maize=M) and put this in the warehouse.

Imagine that warehouses have different rooms for different crops.

The Shopowner's job is to take the crops from this warehouse and make that crops room empty(=N). You have 5 Farmers and 5 Shop Owners

**You need to modify the following C code:**

[https://drive.google.com/file/d/1eVZnBI5oRCCBF9\\_AENWOUH6d4eddQTMV/view?usp=sharing](https://drive.google.com/file/d/1eVZnBI5oRCCBF9_AENWOUH6d4eddQTMV/view?usp=sharing)

### Example Output :

Farmer 3: Insert crops R at 0

Farmer 4: Insert crops W at 1

Farmer 5: Insert crops P at 2

Shop owner 1: Remove crops R from 0  
Shop owner 1: Remove crops W from 1  
Shop owner 1: Remove crops P from 2  
Farmer 4: Insert crops S at 3  
Farmer 4: Insert crops M at 4  
Farmer 4: Insert crops R at 0  
Shop owner 1: Remove crops S from 3  
Shop owner 1: Remove crops M from 4  
ShopOwner1: RNNNN

.....  
.....  
.....  
.....

Shop owner 5: Remove crops W from 1  
Farmer 5: Insert crops P at 2  
Farmer 5: Insert crops S at 3  
Farmer 5: Insert crops M at 4  
Farmer5: NNPSM  
Shop owner 5: Remove crops P from 2  
Shop owner 5: Remove crops S from 3  
Shop owner 5: Remove crops M from 4  
ShopOwner5: NNNNN